

BST 261: Data Science II

Heather Mattie

Department of Biostatistics
Harvard T.H. Chan School of Public Health
Harvard University

March 25, 2019

Linear Algebra Review

- A **scalar** is a single number
- We use $s \in \mathbb{R}$ to denote a real-valued scalar and $n \in \mathbb{N}$ to denote a natural number scalar
- A **vector** is an array of numbers that are arranged in order
- We use $\mathbf{x} \in \mathbb{R}^n$ to denote a vector, where \mathbb{R}^n is the set formed by taking the Cartesian product of \mathbb{R} a total of n times
- Elements of vector \mathbf{x} are identified with x_1, x_2, \dots
- We can think of vectors as identifying points in space, with each element giving the coordinate along a different axis
- All vectors we deal with are column vectors (rather than row vectors) unless we specify otherwise
- We can index a set of elements of a vector by defining a set containing the indices and writing the set as a subscript
- For example we can define $S = \{1, 3, 6\}$ and then write \mathbf{x}_S
- We use the $-$ sign to index the complement of a set, for example, \mathbf{x}_{-S} is the vector containing all elements of \mathbf{x} except for x_1, x_3 , and x_6

- A **matrix** is a 2-D array of numbers
- We use $\mathbf{A} \in \mathbb{R}^{m \times n}$ to denote a real-valued matrix \mathbf{A} of m rows and n columns
- We identify the elements of a matrix using its name and the indices, for example, $A_{1,1}$ is the entry in the first row and first column of \mathbf{A}

$$\mathbf{A} = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix} \quad (1)$$

- We can identify all values of a given index using :
- For example, $\mathbf{A}_{i,:}$ denotes row i of \mathbf{A} and $\mathbf{A}_{:,j}$ denotes column j of \mathbf{A}
- We use subscripts to index matrix-valued expressions, for example, $f(\mathbf{A})_{i,j}$ gives element (i,j) of the matrix computed by applying the function f to \mathbf{A}
- **Tensors** are multidimensional arrays
- We use $A_{i,j,k}$ to identify element (i,j,k) of \mathbf{A}

- Let \mathbf{A} be an $m \times n$ matrix
- The $n \times m$ matrix \mathbf{B} with elements $B_{j,i} = A_{i,j}$ is called the **transpose** of \mathbf{A} and it is denoted by $\mathbf{B} = \mathbf{A}^T$
- Taking the transpose of a matrix amounts to changing rows into columns and vice versa
- The transpose of a matrix is the mirror image of the matrix across a diagonal line, called the **main diagonal**, which runs from the top left element to the bottom right element
- Vectors can be thought of as matrices that contain only one column, and the transpose of a column vector is a row vector and vice versa
- A scalar can be thought of as a matrix with only a single entry, which is its own transpose: $a = a^T$

- Matrices that have the same number of rows and the same number of columns are said to be **conformal** for addition or subtraction
- Matrix addition and subtraction are defined for any two conformal matrices \mathbf{A} and \mathbf{B} , and these operations are performed element-wise: $\mathbf{C} = \mathbf{A} + \mathbf{B}$ where $C_{i,j} = A_{i,j} + B_{i,j}$
- Matrix addition is commutative, that is, $\mathbf{A} + \mathbf{B} = \mathbf{B} + \mathbf{A}$
- A matrix can be multiplied by a scalar by performing that operation on each element of the matrix: $\mathbf{D} = a \cdot \mathbf{B}$ where $D_{i,j} = a \cdot B_{i,j}$
- Additionally some non-conventional shorthand notation is occasionally used:
 - Matrix + scalar: $\mathbf{D} = \mathbf{B} + c$ where $D_{i,j} = B_{i,j} + c$

- The **matrix product** of an $m \times n$ matrix A and an $n \times p$ matrix B is the $m \times p$ matrix C , where the product operation is defined by

$$C_{i,j} = \sum_{k=1}^n A_{i,k} B_{k,j} \quad (2)$$

- We write this as $AB = C$ and refer to it as pre-multiplication of B by A or post-multiplication of A by B
- The **element-wise product**, also known as the **Hadamard product**, of matrices A and B by is denoted by $A \odot B$
- The **dot product** between two vectors x and y of the same dimensionality is the matrix product $x^T y$
- We can therefore think of the matrix product $C = AB$ as computing $C_{i,j}$ as the dot product between row i of A and column j of B

Multiplying Matrices and Vectors

- Matrix multiplication is distributive:

$$A(B + C) = AB + AC \quad (3)$$

- Matrix multiplication is associative:

$$A(BC) = (AB)C \quad (4)$$

- Matrix multiplication is generally not commutative:

$$AB \neq BA \quad (5)$$

- However the dot product between two vectors is commutative:

$$\mathbf{x}^T \mathbf{y} = \mathbf{y}^T \mathbf{x} \quad (6)$$

- The transpose of a matrix product:

$$(AB)^T = B^T A^T \quad (7)$$

- We can write down a system of linear equations:

$$Ax = b \tag{8}$$

- Here $A \in \mathbb{R}^{m \times n}$ is a known matrix, $b \in \mathbb{R}^m$ is a known vector, and $x \in \mathbb{R}^n$ is an unknown vector we would like to solve for
- This can be rewritten as

$$\begin{aligned} A_{1,1}x_1 + A_{1,2}x_2 + \cdots + A_{1,n}x_n &= b_1 \\ A_{2,1}x_1 + A_{2,2}x_2 + \cdots + A_{2,n}x_n &= b_2 \\ &\vdots \\ A_{m,1}x_1 + A_{m,2}x_2 + \cdots + A_{m,n}x_n &= b_m \end{aligned} \tag{9}$$

- **Identity matrix** is a matrix that does not change a vector when multiplied by it
- We denote the identity matrix that preserves n -dimensional vectors as \mathbf{I}_n , where $\mathbf{I} \in \mathbb{R}^{n \times n}$ and

$$\mathbf{I}_n \mathbf{x} = \mathbf{x} \quad \forall \mathbf{x} \in \mathbb{R}^n \quad (10)$$

- In an identity matrix, all entries along the main diagonal are 1, while all the other entries are 0
- The **matrix inverse** of \mathbf{A} is denoted as \mathbf{A}^{-1} and defined as the matrix such that

$$\mathbf{A}^{-1} \mathbf{A} = \mathbf{I}_n \quad (11)$$

- This enables us to solve $\mathbf{A}\mathbf{x} = \mathbf{b}$ as $\mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$

- The inverse \mathbf{A}^{-1} does not always exist, but when it does, several different algorithms can find it in closed form
- The obtained inverse matrix \mathbf{A}^{-1} can then be used to solve the equation $\mathbf{Ax} = \mathbf{b}$ for different values of \mathbf{b}
- In practice, due to the limited precision of a digital computer, the inverse matrix \mathbf{A}^{-1} is rarely used in practice for most software applications
- Instead, algorithms that make use of the value of \mathbf{b} can usually obtain numerically more accurate estimates of \mathbf{x}

- A **linear combination** of some set of vectors $\{\mathbf{v}^{(1)}, \dots, \mathbf{v}^{(n)}\}$ is given by multiplying each vector $\mathbf{v}^{(i)}$ by a corresponding scalar coefficient c_i and adding the results together

$$\sum_i c_i \mathbf{v}^{(i)} \quad (12)$$

- The **span** of a set of vectors $\{\mathbf{v}^{(1)}, \dots, \mathbf{v}^{(n)}\}$ is the set of all points obtainable by their linear combination $\sum_i c_i \mathbf{v}^{(i)}$
- A system of linear equations $\mathbf{A}\mathbf{x} = \mathbf{b}$ has either no solutions or has infinitely many solutions (the only two possibilities) for a given \mathbf{b}
- We can think of the columns of \mathbf{A} as column vectors that specify different directions we can travel in from the origin, and we can think of each element of \mathbf{x} as specifying how far we travel in each of these directions

$$\mathbf{A}\mathbf{x} = \sum_i x_i \mathbf{A}_{:,i} \quad (13)$$

- Determining the number of solutions is then equivalent to determining the number of ways in which we can combine the columns of A to reach b
- Put differently, determining whether $Ax = b$ has a solution amounts to testing whether b is in the span of the columns of A , where the span of the columns of A is known as the **column space** or the **range** of A
- In order for the system $Ax = b$ to have a solution for all values of $b \in \mathbb{R}^m$, we require that the column space of A be all of \mathbb{R}^m
- If any point in \mathbb{R}^m is excluded from the column space, that point is a potential value of b that has no solution
- The requirement that the column space of A be all of \mathbb{R}^m implies immediately that A must have at least m columns
- For example, for a 3×2 matrix A , the target b is in 3 dimensions but x is only in 2 dimensions, so modifying the value of x at best enables us to trace out a 2-D plane within \mathbb{R}^3 ; the equation has a solution if and only if b lies on that plane

Linear Dependence and Span

- A set of vectors is **linearly independent** if no vector in the set is a linear combination of the other vectors
- This means that for the column space of A to encompass all of \mathbb{R}^m , the matrix must contain at least one set of m linearly independent columns, which is both necessary and sufficient for $Ax = b$ to have a solution for every value of b
- The requirement is for exactly m linearly independent columns; no set of m -dimensional vectors can have more than m linearly independent columns, but a matrix with more than m columns may have more than one such set
- For A to have an inverse, we additionally need to ensure that $Ax = b$ has at most one solution for each value of b , meaning that the matrix has at most m columns
- Taken together, this means that the matrix must be square and all the columns must be linearly independent
- A square matrix with linearly dependent columns is called **singular**
- If A is not square, or is square but singular, solving the equation is still possible, but we cannot use matrix inversion to find the solution

- We often need to measure the size of a vector
- This is done using a function called a **norm**, which maps a vector to a non-negative value
- The L^p norm, for $p \in \mathbb{R}, p \geq 1$, is given by

$$\|\mathbf{x}\|_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}} \quad (14)$$

- The L^2 norm is known as the **Euclidean norm**, which is simply the Euclidean distance from the origin to the point identified by \mathbf{x}
- Since this norm is used so frequently, it is sometimes denoted as $\|\mathbf{x}\|$ with the subscript 2 omitted
- The squared L^2 norm is more convenient to work with mathematically and computationally than the L^2 norm itself, and it can be calculated simply as $\mathbf{x}^T \mathbf{x}$

- One other norm that arises is the L^∞ norm, also known as the **max norm**, which simplifies to the absolute value of the element with the largest magnitude in the vector:

$$\|\mathbf{x}\|_\infty = \max_i |x_i| \quad (15)$$

- The most commonly used matrix norm is the **Frobenius norm**

$$\|\mathbf{A}\|_F = \sqrt{\sum_{i,j} A_{i,j}^2} \quad (16)$$

- **Diagonal** matrices consist mostly of zeros and have nonzero entries only along the main diagonal
- We write $\text{diag}(\mathbf{c})$ to denote a square diagonal matrix whose diagonal entries are given by the entries of the vector \mathbf{v}
- Diagonal matrices are of interest in part because multiplying by a diagonal matrix is computationally efficient
- For example, $\text{diag}(\mathbf{v})\mathbf{x} = \mathbf{v} \odot \mathbf{x}$
- The inverse of a square diagonal matrix exists only if every diagonal entry is nonzero, and in that case $\text{diag}(\mathbf{v})^{-1} = \text{diag}([1/v_1, \dots, 1/v_n]^T)$
- Not all diagonal matrices need to be square, i.e., we can have a rectangular diagonal matrix
- Non-square diagonal matrices do not have inverses, but we can still multiply them computationally efficiently

- A **symmetric matrix** is any matrix that is equal to its own transpose: $A = A^T$
- Distance matrices are symmetric because distance functions are symmetric
- A **unit vector** is a vector with **unit norm**:

$$\|x\|_2 = 1 \quad (17)$$

- Vectors x and y are **orthogonal** to each other if $x^T y = 0$
- If both vectors have nonzero norm, this means that they are at a 90 degree angle to each other
- If the vectors are orthogonal and also have unit norm, we call them **orthonormal**

- An **orthogonal matrix** is a square matrix whose rows are mutually orthonormal and whose columns are mutually orthonormal (counterintuitively, not merely orthogonal!)

$$\mathbf{A}^T \mathbf{A} = \mathbf{A} \mathbf{A}^T = \mathbf{I} \quad (18)$$

- This implies that

$$\mathbf{A}^{-1} = \mathbf{A}^T \quad (19)$$

- Inverses of orthogonal matrices are very cheap to compute

- Many mathematical objects can be understood better by breaking them into constituents parts (e.g., decomposing integers into prime factors)
- An **eigenvector** of a square matrix A is a nonzero vector v such that multiplication by A alters only the scale of v :

$$Av = \lambda v \quad (20)$$

- Here the scalar λ is the **eigenvalue** corresponding to eigenvector v
- Since a vector obtained by rescaling v is also an eigenvector, we usually look for unit eigenvectors

- Suppose that a matrix A has n linearly independent eigenvectors $\{\mathbf{v}^{(1)}, \dots, \mathbf{v}^{(n)}\}$ with corresponding eigenvalues $\{\lambda_1, \dots, \lambda_n\}$
- We can concatenate all the eigenvectors to form a matrix V with one eigenvector per column, and we can concatenate the eigenvalues to form a vector λ
- The **eigendecomposition** of A is then given by

$$A = V \text{diag}(\lambda) V^{-1} \quad (21)$$

- Not every matrix can be decomposed into eigenvalues and eigenvectors
- Every real symmetric matrix can be decomposed using only real-valued eigenvectors and eigenvalues

$$A = Q \Lambda Q^T \quad (22)$$

- Here Q is an orthogonal matrix composed of eigenvectors of A and Λ is a diagonal matrix, where the eigenvalue $\Lambda_{i,i}$ is associated with the eigenvector in column i of Q